

On Rust and Categorical Semantics

by Alejandro José Soto Franco

Last updated: July 29, 2025

Hosted on Github · WIP Book

Contents

I High-Level Overview	1
II Preliminaries	4
Definition II.1. Primitive Types in Rust	4
<i>Remark 1</i> (Type Safety as Categorical Discipline)	4
Definition II.2. Functions in Rust	5
<i>Problem 1</i> (<i>Evens and Odds</i>)	5
Definition II.3. Immutability and Mutability	5
Definition II.4. Ownership	6
<i>Remark 2</i> (Ownership as Morphism Restriction)	6
Definition II.5. Borrowing	6
<i>Problem 2</i> (<i>Borrow Checker Conflict</i>)	7
<i>Remark 3</i> (Aliasing Rule and Morphism Uniqueness)	8
Definition II.6. Shadowing	8
III Affine Semantics	10
Definition III.1. Affine Type System	10
<i>Remark 4</i> (Categorical Affine Semantics of Ownership)	10
<i>Problem 3</i> (<i>Ownership vs. Copy</i>)	11
<i>Remark 5</i> (Modal Interpretation)	12

© 2025 Alejandro José Soto Franco

This work is licensed under the
Creative Commons Attribution–NonCommercial–ShareAlike 4.0 International License.

To view a copy of this license, visit:
<http://creativecommons.org/licenses/by-nc-sa/4.0/>

Preface. This is a technical collection of computational and mathematical notes. No claims are made beyond what is written. The work is presented as is. As a work-in-progress draft, what is not yet rigorous may become so later. The author has tried to write only what he can verify.

I High-Level Overview

At its core, Rust enforces a set of invariants corresponding to ownership, borrowing, aliasing, and lifetimes that allow memory-safe, data-race-free programs to be written *without a garbage collector*. Indeed, they are categorical constraints that form a coherent structure over program state. This structure demands a semantics that can:

1. Reflect fine-grained distinctions between move, copy, and borrow,
2. Track the availability and uniqueness of morphisms over program variables,
3. Interpret execution as a homotopically meaningful path through a space of typed states.

Traditional denotational or operational semantics fails to adequately model these distinctions without either overcomplicating the syntax or obscuring the key invariants. Instead, we adopt a semantics based on:

- **Affine type theory:** where variables are resources that may be consumed at most once,
- **Sheaf-theoretic execution models:** where program state is local data with gluing constraints over time,
- **Category-theoretic morphisms:** where ownership and borrow rules translate directly to uniqueness and commutativity of morphisms in a type-preserving category.

This perspective allows us to formalize Rust’s ownership system as a coherent semantic geometry of execution: a system of type-preserving, resource-indexed transformations over program state.

Each of Rust’s core safety rules corresponds to a precise invariant in this execution geometry:

Mutable access	\rightsquigarrow	A unique morphism $f : x \rightarrow x'$ over an affine resource x , disallowing parallel branches,
Shared access	\rightsquigarrow	A family of commuting morphisms $\{f_i : x \rightarrow x\}_{i \in I}$, preserving read-only semantics,
Lifetimes	\rightsquigarrow	Restriction of morphism support to a region $U \subseteq \mathcal{T}$, where \mathcal{T} is program time,
Move semantics	\rightsquigarrow	Affine context contraction $\Gamma \vdash e : \tau \implies \Gamma \setminus \{x\} \vdash e' : \tau'$, consuming x .

By treating these semantics as the *primary object of study*, we are able to:

- Define well-formed execution paths as categorical compositions,
- Make precise statements about soundness, race-freedom, and memory correctness,
- Generalize Rust’s ideas to new safe languages or embedded DSLs with provable safety guarantees.

Ultimately, Rust offers a bridge between high-performance computation and formal reasoning. To cross it rigorously, we must treat the semantics not as a byproduct, but as a guiding structure. The goal of this document is to expose that structure — to show that Rust is not merely a language with rules, but a *geometry of execution* governed by types, and to construct the categorical semantics that makes this geometry coherent, composable, and verifiable.

The following are tables roughly corresponding to the topics discussed in each section:

Core Rust Constructs and Affine Semantics

Rust Construct	Semantic Interpretation
<code>let x = 5;</code>	Immutable binding. Disallows duplication (affine logic).
<code>let mut x = 5;</code>	Mutable binding. Enables linear update; weakening allowed.
<code>fn f(x: T) -> U</code>	Morphism $f : T \rightarrow U$ in affine type category.
<code>let b = a;</code>	Ownership move. Consumes 'a' unless 'T: Copy'.
<code>&a</code>	Shared borrow. Modal access $\Box T$; comonadic read-only.
<code>&mut a</code>	Exclusive borrow. Linear access $\triangleright T$; no aliasing.

Structural Rules of Affine Logic

Rule	Affine Logic	Rust Behavior
Weakening	Unused bindings permitted	Declared vars may go unused
Contraction	Disallowed	No duplication unless Copy
Exchange	Permitted	Statement order flexible

Ownership, Copy, and Borrowing

Operation	Condition	Semantic Class
Move	Default	Consumes value (linear logic)
Copy	If T: Copy	Duplication allowed; contraction admitted
Immutable borrow	Always	Comonadic: $\Box T$
Mutable borrow	Exclusive	Linear ephemeral: $\triangleright T$
Shadowing	Type-safe	New binding masks prior

Categorical Interpretation of Rust

Rust Concept	Categorical Structure
Types T	Objects in affine symmetric monoidal category AffType
Functions fn	Morphisms $f : T \rightarrow U$
Ownership Move	Consuming morphism (non-duplicative)
Copy	Identity-preserving duplication morphism
Borrowing	Modal morphisms (\Box, \triangleright)
Traits	Predicate subobjects or classifier constraints

Borrow Checker as Categorical Constraint

Invariant	Categorical or Logical Justification
No <code>&mut</code> if <code>&</code> exists	Morphism uniqueness in affine category
Only one <code>&mut</code> at a time	Linear typing discipline
Lifetimes	Indexed morphisms over regions; natural transformations
Reborrow	Subobject relation under scope restriction

Trait Semantics and Resource Flow

Trait	Purpose	Semantic Role
Copy	Implicit duplication	Contraction permitted
Clone	Explicit duplication	Manual morphism invocation
Send	Thread-safe transfer	Safe morphism across threads
Sync	Shared thread access	Shared morphism closure
Drop	Cleanup on scope end	Terminal morphism (finalizer)

Modal Typing Summary

Symbol	Rust Meaning	Typing Semantics
$\Box T$	<code>&T</code>	Comonadic, shared read-only
$\triangleright T$	<code>&mut T</code>	Linear, exclusive use
T	Owned value	Affine object (non-duplicable)

Work-in-Progress Draft
Alejandro Soto Franco, CC BY-NC-SA

II Preliminaries

This section introduces core syntactic and semantic elements of the **Rust** programming language that will recur throughout our mathematical treatment. We begin with primitive types, basic function definition, and ownership semantics. These elements will later be given categorical and geometric interpretations.

Definition II.1. Primitive Types in Rust.

The set of primitive types in **Rust** includes the fixed-size, compile-time known scalar types listed below. These types are stratified by logical, integer, and floating-point structure, each with a specified bit-width and memory layout.

Type	Bits	Signedness	Structure
<code>bool</code>	1	n/a	Logical type
<code>char</code>	32	Unicode scalar	Character code point
<code>u8</code>	8	Unsigned	Integer
<code>u16</code>	16	Unsigned	Integer
<code>u32</code>	32	Unsigned	Integer
<code>u64</code>	64	Unsigned	Integer
<code>usize</code>	Platform-dependent	Unsigned	Pointer-sized Integer
<code>i8</code>	8	Signed	Integer
<code>i16</code>	16	Signed	Integer
<code>i32</code>	32	Signed	Integer
<code>i64</code>	64	Signed	Integer
<code>isize</code>	Platform-dependent	Signed	Pointer-sized Integer
<code>f32</code>	32	IEEE 754	Floating-point
<code>f64</code>	64	IEEE 754	Floating-point

Types such as `u32`, `i64`, and `f64` form the basis of numeric computation, with arithmetic operators and ordering defined over them. The type `bool` encodes the two-valued logical space $\{\text{true}, \text{false}\}$, and `char` denotes valid Unicode scalar values, not necessarily ASCII. The types `usize` and `isize` track platform-dependent word width and are used for memory indexing. \lrcorner

Remark 1. (Type Safety as Categorical Discipline)

The type system of **Rust** is nominal, statically checked, and strongly typed. Every value in a well-formed program is assigned a single, unambiguous type that is fully determined at compile time. Type annotations, inference, and trait resolution collectively ensure that no value can be interpreted at runtime in a way inconsistent with its static type.

From a semantic perspective, this corresponds to a categorical discipline in which:

- *Types* are treated as objects in a category $\mathcal{C}_{\text{Rust}}$,
- *Expressions* correspond to morphisms $f : \tau_1 \rightarrow \tau_2$,
- *Well-typedness* enforces that all morphisms are defined only on valid objects, with respect to ownership, borrowing, and lifetimes,

- *Type errors* reflect *undefined morphisms*, and are excluded from the semantics entirely.

Thus, the typing judgment

$$\Gamma \vdash e : \tau$$

may be interpreted as asserting the existence of a partial morphism

$$f_e : \text{Dom}_\Gamma \dashrightarrow \tau$$

defined only under the constraints enforced by the context Γ . This partiality is not due to non-termination or dynamic checks, but arises from the *logical geometry* of affine resources and the static structure of the program.

Categorically, the restriction to well-typed programs corresponds to working in a *typed subcategory* $\mathcal{C}_{\text{safe}} \subseteq \mathcal{C}_{\text{Rust}}$, where morphisms preserve not only type structure but also aliasing discipline and memory soundness. As such, type safety in **Rust** is not merely a compiler check — it is the formal criterion by which the composability and realizability of morphisms in the execution space is guaranteed. We will later see that lifetimes impose sheaf-theoretic support conditions, and ownership corresponds to affine contraction laws, all of which are enforced through this categorical structure. \rightarrow

Definition II.2. Functions in Rust.

A function is declared via the syntax:

```
fn f(x: i32) -> i32 {
    x * x
}
```

This defines a mapping $f : \mathbb{Z} \rightarrow \mathbb{Z}$ via the rule $x \mapsto x^2$. The compiler checks that all branches return the declared type. \lrcorner

Problem 1. (Evens and Odds)

Write a function in **Rust** that determines whether an integer is even. \lrcorner

Solution.

```
fn is_even(n: i32) -> bool {
    n % 2 == 0
}
```

This uses modulo arithmetic and returns a **bool**.

Definition II.3. Immutability and Mutability.

In **Rust**, variables are immutable by default:

```
let x = 5;           // x is immutable
let mut y = 5;       // y is mutable
```

This immutability reflects an affine logic perspective. In later sections we will interpret this as the default denial of contraction. \lrcorner

Example. In contrast to many languages, `Rust` enforces that mutation must be explicit. The ownership and borrowing rules ensure safe concurrency and eliminate data races by construction. \diamond

Definition II.4. Ownership.

Each value in `Rust` has a single owner at any time. When ownership is transferred (moved), the previous binding becomes invalid:

```
let s1 = String::from("hello");
let s2 = s1; // ownership moved to s2
// s1 is no longer valid
```

This model ensures no two variables alias the same memory without the borrow checker knowing. \lrcorner

Remark 2. (Ownership as Morphism Restriction)

`Rust`'s ownership system enforces a precise constraint on how program state evolves: each value has a unique owner, and this ownership must be explicitly transferred, borrowed, or consumed. From a semantic perspective, this corresponds to a structural constraint that determines which morphisms are defined between typed program states, with execution valid only when transitions preserve ownership.

Let \mathcal{C} be a category whose objects are well-typed program states, and whose morphisms represent valid transitions between these states induced by program expressions. In this setting, *ownership* corresponds to a condition on morphism formation: a morphism $f : x \rightarrow y$ is defined only when the value x has not been aliased, borrowed mutably, or consumed elsewhere in the execution trace. The morphism f must *consume* x from the domain, reflecting that the resource is no longer available thereafter.

Formally, we express this as a *partial morphism restriction*:

$$\text{Hom}_{\text{own}}(x, y) \subseteq \text{Hom}_{\mathcal{C}}(x, y)$$

where $\text{Hom}_{\text{own}}(x, y)$ denotes morphisms respecting ownership constraints. This substructure of the full morphism space captures the affine typing discipline: each resource (variable binding) may be used at most once in constructing morphisms, unless the type supports duplication (e.g., via 'Copy').

This restriction induces a geometric structure on the manifold of program execution: not all paths between states are valid, only those corresponding to ownership-preserving morphisms. Executing a program then becomes equivalent to tracing a directed path through this manifold, where each step consumes inputs in accordance with affine type constraints and produces new, uniquely owned outputs.

In later sections, we will extend this picture to include borrow semantics as morphisms into read-only fibers, and lifetimes as region-restricted morphism supports. Ownership, in this unified model, is not a language feature but a manifestation of a deeper constraint: *that morphisms over state must respect the affine geometry of consumption and exclusivity.* \rightarrow

Definition II.5. Borrowing.

`Rust` allows references to values without transferring ownership:

```
fn len(s: &String) -> usize {
    s.len()
}
```


}

Here `&String` is an immutable reference. Mutable references `&mut T` are exclusive and non-aliasable. ┘

Problem 2. (Borrow Checker Conflict)

Explain why the following Rust code results in a compile-time error:

```
let mut x = 5;
let y = &x;
let z = &mut x;
```

┘

Solution. The error is due to a violation of Rust’s **exclusive aliasing invariant**, a core enforcement of its **affine type system**. This rule is rooted in affine logic, which permits weakening (discarding unused resources) but disallows contraction (duplicating them). Rust translates this into the following runtime constraint:

- You may have either one mutable reference (`&mut T`) **or**
- Any number of immutable references (`&T`),
- **But not both simultaneously.**

In the given code:

- `x` is mutable and is first immutably borrowed as `y : &i32`,
- Then `z : &mut i32` is created while `y` is still in scope,

violating Rust’s borrowing rules. From the perspective of type systems:

- Immutable borrows preserve resource identity and do not consume it.
- Mutable borrows **consume** the ability to alias or share, requiring exclusive access.

In categorical semantics, this behavior can be modeled in an affine symmetric monoidal category:

- Let $\mathbf{AffType}$ be a category of affine types,
- A value of type T is an object,
- A mutable borrow is a morphism $\mathbf{mut_ref}_T : T \rightarrow \&\mathbf{mut} T$ that consumes T ,
- An immutable borrow is a morphism $\mathbf{shared_ref}_T : T \rightarrow \Box T$ where \Box is a comonadic modality preserving read-only access.

The conflict arises because you attempt to create `z` while `y` still exists. This would force a diagram with both `&T` and `&mut T` arrows in the same scope, violating uniqueness of morphisms in the affine category. **Thus, the compiler prevents this construct at compile-time to uphold memory safety and preserve the semantics of linear resource flow.**

Remark 3. (Aliasing Rule and Morphism Uniqueness)

This restriction prevents data races and enables non-interleaved reasoning about program state. From a semantics perspective, we will later encode this as a uniqueness condition on morphisms.

To formalize this, recall that in Rust, at any given time, one of the following must hold for a variable binding:

- (i) Exactly one mutable reference exists, and no immutable references,
- (ii) Any number of immutable references exist, and no mutable references.

We interpret the program heap (or semantic memory) as a fibered category of stateful objects. Let \mathcal{C} be the category of memory configurations, and let each morphism $f : x \rightarrow y$ represent a permitted transformation of state via a function or method application.

Then, the *aliasing rule* corresponds to the requirement that:

$$\forall x \in \text{Ob}(\mathcal{C}), \quad \exists! f_x^{\text{mut}} : x \rightarrow x' \quad \text{or} \quad \forall i, f_i^{\text{imm}} : x \rightarrow x \quad \text{such that} \quad [f_i, f_j] = 0 \quad \forall i, j. \quad (1)$$

That is, either:

- There exists a unique morphism representing mutable access (exclusive right to evolve the object), or
- There exists a commuting family of morphisms representing immutable access (read-only projections).

This ensures the state evolution diagram remains non-branching under mutation and commutative under observation:

$$\begin{array}{ccc} & x & \\ f_i^{\text{imm}} \swarrow & & \searrow f_j^{\text{imm}} \\ x & \xrightarrow{f^{\text{mut}}} & x' \\ & \xlongequal{\quad} & x \end{array}$$

Categorically, this induces a *uniqueness constraint* on the morphisms permitted within a given context. We can write this as a context morphism condition:

$$\Gamma \vdash \exists! f : x \rightarrow y \quad \text{or} \quad \Gamma \vdash \{f_i : x \rightarrow x\}_{i \in I} \text{ with } [f_i, f_j] = 0. \quad (2)$$

This encoding provides a semantic justification for Rust's aliasing rules, and also serves as a foundational rule for safe, concurrent execution. Uniqueness of morphisms corresponds to thread safety guarantees and compositional determinism.

This rule will later appear as a coherence law in our Ω -type model of Rust's ownership system as part of the categorical semantics framework we seek to build across all possible typed geometric systems, including Clifford-algebraic, cohomological, and modal resource-sensitive logics. \rightarrow

Definition II.6. Shadowing.

Rust allows a variable to be re-bound with the same name:

```
let x = 5;
let x = x + 1;
```

The second `x` shadows the first. This differs from mutation and does not require `mut`. \lrcorner

Example. Shadowing is useful for type transformations:

```
let spaces = "   ";  
let spaces = spaces.len(); // spaces is now a usize
```

i.e., a `&str` becomes a `usize`.

◇

Proposition II.7. Shadowing's Usefulness.

Shadowing introduces a new binding in the same scope, which behaves like a fresh immutable variable. It does *not* modify the original binding, and both bindings may coexist conceptually during compilation, though only the latest is accessible.

Formally, if `let x = e;` and then `let x = f(x);`, the second `x` is independent of the memory or mutability status of the first.

Therefore, shadowing preserves immutability and supports safe re-binding semantics without aliasing or lifetime issues.

┘

Work-in-Progress Draft
Alejandro Soto Franco, CC BY-NC-SA

III Affine Semantics

Definition III.1. Affine Type System.

An *affine type system* is a typing discipline in which each bound variable may be used at most once. Formally, given a context Γ , the judgment $\Gamma \vdash e : \tau$ implies that every variable in Γ is consumed at most once during the evaluation of e . \lrcorner

Remark 4. (Categorical Affine Semantics of Ownership)

Rust's ownership and borrowing system can be rigorously modeled using typed categorical semantics:

- The type system is *affine*, meaning variables must be used at most once. This disallows contraction but allows weakening.
- Affine logic forbids the structural duplication of assumptions:

$$(\text{Contr}) \quad \frac{\Gamma, x : A, x : A \vdash e : B}{\Gamma, x : A \vdash e : B}$$

but permits their omission:

$$(\text{Weak}) \quad \frac{\Gamma \vdash e : B}{\Gamma, x : A \vdash e : B}$$

- Categorically, Rust's semantics can be embedded in a *typed symmetric monoidal category* **AffType**:
 - Objects are types equipped with resource annotations (e.g., $x : T$),
 - Morphisms $f : A \rightarrow B$ represent computations consuming A to produce B without duplication,
 - The tensor product $A \otimes B$ models concurrent resource availability.
- Borrowing introduces modal structure:
 - Shared borrows ' $\&T$ ' admit a comonadic modality \Box , enabling read-only access,
 - Exclusive borrows ' $\&\text{mut } T$ ' model ephemeral, linear use of the resource without aliasing.

\rightrightarrows

Theorem III.2. No Contraction in Affine Semantics.

Let $\Gamma, x : \tau \vdash e : \sigma$ be a valid typing judgment in an affine type system. Then there does not exist a derivation rule of the form:

$$\frac{\Gamma, x_1 : \tau, x_2 : \tau \vdash e[x/x_1, x/x_2] : \sigma}{\Gamma, x : \tau \vdash e[x/x_1, x/x_2] : \sigma} \quad (\text{Contraction})$$

unless the type τ satisfies $\tau \in \text{Copy}$, where **Copy** is the set of types that admit duplicable semantics under Rust's trait system. \lrcorner

Proof. In an affine type system, variable bindings are treated as linear resources: each bound variable must be used at most once in the evaluation of an expression. The contraction rule, which allows a variable $x : \tau$ to be duplicated into multiple instances x_1, x_2 within the same typing context, violates this linear usage constraint.

By definition, contraction asserts the admissibility of a structural rule of the form:

$$(\text{Contr}) \quad \Gamma, x : \tau \vdash e[x/x_1, x/x_2] : \sigma \quad \text{where } x_1, x_2 : \tau$$

which would allow two uses of x where only one binding exists.

Such a rule is not derivable in affine systems. In Rust, this semantic restriction is enforced at the type level: only types implementing the `Copy` trait may be duplicated implicitly. For non-`Copy` types, such as `Vec<T>`, `String`, or user-defined types with heap-allocated state, duplication without explicit cloning constitutes a violation of memory safety and is therefore disallowed.

Thus, in Rust's affine model, the absence of contraction ensures that each instance of a value is used exactly once unless its type is explicitly duplicable. This enforces exclusive ownership and prevents aliasing, double frees, and race conditions. The uniqueness of usage is preserved categorically by modeling each variable-to-expression relationship as a *partial morphism* with strict consumption semantics.

Therefore, contraction is ruled out for general τ , and only recoverable when τ admits duplication under an auxiliary semantic judgment:

$$\tau \in \text{Copy} \quad \Rightarrow \quad \text{Contraction admissible.}$$

□

Problem 3. (Ownership vs. Copy)

In Rust, the following code is valid:

```
let x = 3;           // i32 implements Copy
let y = x;
let z = x;           // x still valid
```

But the following fails:

```
let s1 = String::from("hello");
let s2 = s1;
let s3 = s1;         // error: use of moved value
```

Explain the underlying affine typing reasons. ┐

Solution. The type `i32` implements the `Copy` trait, so assignment creates a bitwise duplicate without violating affine semantics. In contrast, `String` owns heap-allocated data and does not implement `Copy`; thus, moving it transfers ownership. Reusing `s1` after the move violates the single-use constraint of affine logic.

Proposition III.3. Borrowing as Modal Operation.

Let $\Gamma \vdash e : \tau$ be a well-typed expression in an affine type system. Then forming a shared reference `&e` corresponds to a modal operation in the type system:

$$\frac{\Gamma \vdash e : \tau}{\Gamma \vdash \&e : \Diamond \tau}$$

where $\Diamond\tau$ denotes a type under a modality that grants permission to observe e without consuming it.

Properties:

1. **Non-consumption:** The original term e remains valid in the context Γ . Borrowing does not move or invalidate ownership.
2. **Duplication:** Values of type $\Diamond\tau$ may be duplicated freely. That is, $\Diamond\tau \vdash \Diamond\tau \otimes \Diamond\tau$.
3. **Read-only Semantics:** No mutation is permitted through $\Diamond\tau$; access is observational only.
4. **Modal Interpretation:** The operator \Diamond behaves like a comonadic modality from linear logic, reflecting duplicability and non-consumption.

Therefore, shared borrowing in Rust aligns with a modal typing rule, where $\Diamond\tau$ encapsulates safe, aliasable observation without affecting the linear flow of ownership. \lrcorner

Remark 5. (Modal Interpretation)

Borrowing introduces a modality into the type system that distinguishes between consuming and non-consuming access.

Shared borrow: $\&e$ corresponds to a type $\Diamond\tau$, where:

- Observation is permitted,
- Duplication is permitted,
- Mutation is prohibited.

Mutable borrow: $\&\text{mut } e$ introduces a separate, *affine-exclusive* modality $\blacklozenge\tau$ with the properties:

- Observation is permitted,
- Mutation is permitted,
- Duplication is prohibited.

These modalities reflect a structural refinement of access rights:

$$\text{Ownership} \Rightarrow \text{Affine Modalities} \Rightarrow \text{Usage Permissions}$$

and clarify the design of Rust's borrow checker as enforcing modal constraints over a linear-affine core. \rightarrow